

Physics-guided characterisation-time reduction for second-life LFP cells:

how few cycles does PyBaMM's reaction-limited SEI need?

Krishna Chaitanya Vaddepally[†]

[†]Presenting author. Turno (Blubble Pvt. Ltd.), India.


battery.product@turno.club

Category: Physics-based Models • **Preference:** Oral (poster acceptable)


Motivation. Second-life repurposing of used lithium-iron-phosphate (LFP) cells for stationary energy storage is bottlenecked by characterisation cost: pack builders must estimate remaining useful life (RUL) with no beginning-of-life data and limited laboratory cycling time. We ask: *how many measured cycles does a physics-based degradation model need before it can predict the remaining life of a used cell?*

Method. We calibrate PyBaMM's reaction-limited SEI degradation model (Prada2013 chemistry, isothermal 25 °C) against the measured fade slope of a used cell using only its first K cycles, then run the calibrated model to the full measured horizon and score against the held-out cycles $K-N$. A single scalar (SEI exchange current density j_{SEI}) is fit per cell by bisection on $\log_{10}(j_{SEI})$; per-cell open-circuit voltage, initial state-of-health, and ohmic resistance are constrained from HPPC and OCV-SoC characterisation. We sweep $K \in \{50, 100, 200, 400, 800\}$ across seven MFR_A-cohort cells with clean long-trajectory records ($N > 1000$ cycles each) that pass an upstream reaction-limited-shape filter.

Result. Prediction RMSE on held-out cycles drops from a median of 11.2 pp SoH at $K = 50$ to 1.4 pp at $K = 100$ and 0.5 pp at $K = 800$ (Fig. 1a). At $K = 400$, all 7/7 cells achieve $RMSE_{test} < 3$ pp SoH: **400 vs. ≈ 1200 measured cycles per cell ($3\times$ reduction, ≈ 4 months of test-lab time saved at 0.25 C).** Five of seven cells already converge at $K = 100$ with end-of-trajectory error below 0.4 pp (Fig. 1b) — a **$12\times$ reduction** for that sub-cohort. The remaining two exhibit a post-formation recovery transient (positive apparent slope across cycles 50–200) and require $K \geq 400$; this signal itself classifies cells for correct calibration windowing.



(a) Held-out RMSE vs. training window K ; red dashed = 3 pp target.



(b) $K=100$ -cycle fit predicts remaining 1345 cycles of cell 07 to 0.29 pp SoH RMSE.

Scope & PyBaMM usage. The result holds only within the reaction-limited SEI regime; cells where graphite loss-of-active-material dominates (identified upstream by non-monotonic-early or three-regime shape) fall outside this envelope because PyBaMM's stress-driven LAM formulation lacks saturation and overfits short calibration windows. PyBaMM (v26.4.3) is the calibration target and extrapolation engine; per-cell characterisation is mapped into the Prada2013 parameter set via `build_pybamm_parameters`, and bisection over $\log_{10}(j_{SEI})$ uses cached 10-cycle PyBaMM runs. Scripts released at <https://github.com/VKC-Turno/PyBaMM-Neural-ODE-RUL>.

Acknowledgements. The author thanks Sunil Kumar Rawat (Imperial College London) for developing the initial cell parameter identification and first-life PyBaMM degradation model on which this work builds.

References. [1] V. Sulzer et al., *J. Open Res. Softw.* 9(1) 14–21 (2021). [2] M. Prada et al., *J. Electrochem. Soc.* 160(4) A616–A628 (2013).